

Multi_Layer_Perceptron

April 8, 2020

```
[0]: import numpy as np
import tensorflow as tf
from tensorflow.keras.layers import Input, Flatten, Dense
from tensorflow.keras.models import Model
from tensorflow.keras.datasets import mnist
from matplotlib import pyplot as plt
```

```
[0]: (X_train, y_train), (X_test, y_test) = mnist.load_data()

X_train = X_train / 255.
X_test = X_test / 255.
```

```
[25]: input_img = Input(shape=(28, 28))
x = Flatten()(input_img)
x = Dense(50, activation="relu")(x)
classification_layer = Dense(10)(x)

model = Model(input_img, classification_layer)
model.compile(optimizer=tf.keras.optimizers.Adam(lr=0.01),
              loss=tf.keras.losses.
↳SparseCategoricalCrossentropy(from_logits=True),
              metrics=["accuracy"])
model.summary()

model.fit(X_train, y_train,
         epochs=10,
         batch_size=128,
         shuffle=True,
         validation_data=(X_test, y_test))
```

Model: "model_7"

Layer (type)	Output Shape	Param #
input_10 (InputLayer)	[(None, 28, 28)]	0
flatten_8 (Flatten)	(None, 784)	0

dense_14 (Dense)	(None, 50)	39250
------------------	------------	-------

dense_15 (Dense)	(None, 10)	510
------------------	------------	-----

```

Total params: 39,760
Trainable params: 39,760
Non-trainable params: 0

```

```

Epoch 1/10
469/469 [=====] - 1s 3ms/step - loss: 0.2623 -
accuracy: 0.9216 - val_loss: 0.1688 - val_accuracy: 0.9490
Epoch 2/10
469/469 [=====] - 1s 3ms/step - loss: 0.1392 -
accuracy: 0.9583 - val_loss: 0.1367 - val_accuracy: 0.9586
Epoch 3/10
469/469 [=====] - 1s 3ms/step - loss: 0.1114 -
accuracy: 0.9655 - val_loss: 0.1449 - val_accuracy: 0.9589
Epoch 4/10
469/469 [=====] - 1s 3ms/step - loss: 0.1016 -
accuracy: 0.9697 - val_loss: 0.1468 - val_accuracy: 0.9581
Epoch 5/10
469/469 [=====] - 1s 3ms/step - loss: 0.0863 -
accuracy: 0.9738 - val_loss: 0.1397 - val_accuracy: 0.9640
Epoch 6/10
469/469 [=====] - 1s 3ms/step - loss: 0.0823 -
accuracy: 0.9746 - val_loss: 0.1618 - val_accuracy: 0.9565
Epoch 7/10
469/469 [=====] - 1s 3ms/step - loss: 0.0758 -
accuracy: 0.9768 - val_loss: 0.1428 - val_accuracy: 0.9653
Epoch 8/10
469/469 [=====] - 1s 3ms/step - loss: 0.0740 -
accuracy: 0.9772 - val_loss: 0.1493 - val_accuracy: 0.9660
Epoch 9/10
469/469 [=====] - 1s 3ms/step - loss: 0.0689 -
accuracy: 0.9785 - val_loss: 0.1400 - val_accuracy: 0.9685
Epoch 10/10
469/469 [=====] - 1s 3ms/step - loss: 0.0632 -
accuracy: 0.9803 - val_loss: 0.1485 - val_accuracy: 0.9674

```

[25]: <tensorflow.python.keras.callbacks.History at 0x7f001f298780>

```

[26]: input_img = Input(shape=(28, 28))
      x = Flatten()(input_img)
      x = Dense(150, activation="relu")(x)
      classification_layer = Dense(10)(x)

      model = Model(input_img, classification_layer)

```

```

model.compile(optimizer=tf.keras.optimizers.Adam(lr=0.01),
              loss=tf.keras.losses.
↳SparseCategoricalCrossentropy(from_logits=True),
              metrics=["accuracy"])
model.summary()

model.fit(X_train, y_train,
        epochs=10,
        batch_size=128,
        shuffle=True,
        validation_data=(X_test, y_test))

```

Model: "model_8"

Layer (type)	Output Shape	Param #
input_11 (InputLayer)	[(None, 28, 28)]	0
flatten_9 (Flatten)	(None, 784)	0
dense_16 (Dense)	(None, 150)	117750
dense_17 (Dense)	(None, 10)	1510

Total params: 119,260
 Trainable params: 119,260
 Non-trainable params: 0

```

Epoch 1/10
469/469 [=====] - 2s 4ms/step - loss: 0.2224 -
accuracy: 0.9329 - val_loss: 0.1084 - val_accuracy: 0.9671
Epoch 2/10
469/469 [=====] - 2s 4ms/step - loss: 0.1077 -
accuracy: 0.9672 - val_loss: 0.1129 - val_accuracy: 0.9653
Epoch 3/10
469/469 [=====] - 2s 4ms/step - loss: 0.0851 -
accuracy: 0.9737 - val_loss: 0.1306 - val_accuracy: 0.9633
Epoch 4/10
469/469 [=====] - 2s 4ms/step - loss: 0.0800 -
accuracy: 0.9750 - val_loss: 0.1037 - val_accuracy: 0.9716
Epoch 5/10
469/469 [=====] - 2s 4ms/step - loss: 0.0647 -
accuracy: 0.9801 - val_loss: 0.1398 - val_accuracy: 0.9668
Epoch 6/10
469/469 [=====] - 2s 4ms/step - loss: 0.0643 -
accuracy: 0.9808 - val_loss: 0.1442 - val_accuracy: 0.9684
Epoch 7/10

```

```

469/469 [=====] - 2s 4ms/step - loss: 0.0634 -
accuracy: 0.9814 - val_loss: 0.1412 - val_accuracy: 0.9683
Epoch 8/10
469/469 [=====] - 2s 4ms/step - loss: 0.0615 -
accuracy: 0.9824 - val_loss: 0.1272 - val_accuracy: 0.9717
Epoch 9/10
469/469 [=====] - 2s 4ms/step - loss: 0.0477 -
accuracy: 0.9862 - val_loss: 0.1495 - val_accuracy: 0.9723
Epoch 10/10
469/469 [=====] - 2s 4ms/step - loss: 0.0519 -
accuracy: 0.9856 - val_loss: 0.1493 - val_accuracy: 0.9702

```

[26]: <tensorflow.python.keras.callbacks.History at 0x7f001f35ef60>

```

[35]: input_img = Input(shape=(28, 28))
x = Flatten()(input_img)
x = Dense(50, activation="relu", kernel_regularizer=tf.keras.regularizers.
↳l2(l=0.001))(x)
classification_layer = Dense(10)(x)

model = Model(input_img, classification_layer)
model.compile(optimizer=tf.keras.optimizers.Adam(lr=0.01),
              loss=tf.keras.losses.
↳SparseCategoricalCrossentropy(from_logits=True),
              metrics=["accuracy"])
model.summary()

model.fit(X_train, y_train,
        epochs=10,
        batch_size=128,
        shuffle=True,
        validation_data=(X_test, y_test))

```

Model: "model_17"

Layer (type)	Output Shape	Param #
input_20 (InputLayer)	[(None, 28, 28)]	0
flatten_18 (Flatten)	(None, 784)	0
dense_34 (Dense)	(None, 50)	39250
dense_35 (Dense)	(None, 10)	510

```

Total params: 39,760
Trainable params: 39,760
Non-trainable params: 0

```

```

-----
Epoch 1/10
469/469 [=====] - 1s 3ms/step - loss: 0.4018 -
accuracy: 0.9128 - val_loss: 0.3123 - val_accuracy: 0.9390
Epoch 2/10
469/469 [=====] - 1s 3ms/step - loss: 0.3264 -
accuracy: 0.9380 - val_loss: 0.3508 - val_accuracy: 0.9288
Epoch 3/10
469/469 [=====] - 1s 3ms/step - loss: 0.3132 -
accuracy: 0.9417 - val_loss: 0.3090 - val_accuracy: 0.9392
Epoch 4/10
469/469 [=====] - 1s 3ms/step - loss: 0.3144 -
accuracy: 0.9430 - val_loss: 0.2880 - val_accuracy: 0.9468
Epoch 5/10
469/469 [=====] - 1s 3ms/step - loss: 0.3057 -
accuracy: 0.9450 - val_loss: 0.3624 - val_accuracy: 0.9268
Epoch 6/10
469/469 [=====] - 1s 3ms/step - loss: 0.3090 -
accuracy: 0.9431 - val_loss: 0.2758 - val_accuracy: 0.9542
Epoch 7/10
469/469 [=====] - 1s 3ms/step - loss: 0.3030 -
accuracy: 0.9452 - val_loss: 0.3045 - val_accuracy: 0.9436
Epoch 8/10
469/469 [=====] - 1s 3ms/step - loss: 0.2968 -
accuracy: 0.9463 - val_loss: 0.2992 - val_accuracy: 0.9449
Epoch 9/10
469/469 [=====] - 1s 3ms/step - loss: 0.2971 -
accuracy: 0.9452 - val_loss: 0.3073 - val_accuracy: 0.9428
Epoch 10/10
469/469 [=====] - 1s 3ms/step - loss: 0.2899 -
accuracy: 0.9469 - val_loss: 0.2822 - val_accuracy: 0.9498

```

[35]: <tensorflow.python.keras.callbacks.History at 0x7f0017f90390>

```

[36]: input_img = Input(shape=(28, 28))
x = Flatten()(input_img)
x = Dense(250, activation="relu", kernel_regularizer=tf.keras.regularizers.
↳l2(l=0.001))(x)
classification_layer = Dense(10)(x)

model = Model(input_img, classification_layer)
model.compile(optimizer=tf.keras.optimizers.Adam(lr=0.01),
              loss=tf.keras.losses.
↳SparseCategoricalCrossentropy(from_logits=True),
              metrics=["accuracy"])
model.summary()

```

```
model.fit(X_train, y_train,
          epochs=10,
          batch_size=128,
          shuffle=True,
          validation_data=(X_test, y_test))
```

Model: "model_18"

Layer (type)	Output Shape	Param #
input_21 (InputLayer)	[(None, 28, 28)]	0
flatten_19 (Flatten)	(None, 784)	0
dense_36 (Dense)	(None, 250)	196250
dense_37 (Dense)	(None, 10)	2510

Total params: 198,760

Trainable params: 198,760

Non-trainable params: 0

Epoch 1/10

469/469 [=====] - 2s 5ms/step - loss: 0.4713 - accuracy: 0.9169 - val_loss: 0.4151 - val_accuracy: 0.9310

Epoch 2/10

469/469 [=====] - 2s 5ms/step - loss: 0.3876 - accuracy: 0.9366 - val_loss: 0.3628 - val_accuracy: 0.9380

Epoch 3/10

469/469 [=====] - 2s 5ms/step - loss: 0.3545 - accuracy: 0.9412 - val_loss: 0.3443 - val_accuracy: 0.9449

Epoch 4/10

469/469 [=====] - 2s 5ms/step - loss: 0.3344 - accuracy: 0.9439 - val_loss: 0.3107 - val_accuracy: 0.9479

Epoch 5/10

469/469 [=====] - 2s 5ms/step - loss: 0.3278 - accuracy: 0.9437 - val_loss: 0.3043 - val_accuracy: 0.9489

Epoch 6/10

469/469 [=====] - 2s 5ms/step - loss: 0.3036 - accuracy: 0.9481 - val_loss: 0.3055 - val_accuracy: 0.9406

Epoch 7/10

469/469 [=====] - 2s 5ms/step - loss: 0.3078 - accuracy: 0.9466 - val_loss: 0.3018 - val_accuracy: 0.9443

Epoch 8/10

469/469 [=====] - 2s 5ms/step - loss: 0.3102 - accuracy: 0.9456 - val_loss: 0.2688 - val_accuracy: 0.9556

Epoch 9/10

```

469/469 [=====] - 2s 5ms/step - loss: 0.2951 -
accuracy: 0.9484 - val_loss: 0.3266 - val_accuracy: 0.9393
Epoch 10/10
469/469 [=====] - 2s 5ms/step - loss: 0.3018 -
accuracy: 0.9465 - val_loss: 0.2969 - val_accuracy: 0.9453

```

[36]: <tensorflow.python.keras.callbacks.History at 0x7f00183259b0>

	Classification accuracy	
	Training	Testing
50HLN+no regularization	0.9803	0.9674
50HLN+L2 regularization	0.9469	0.9498
250HLN+no regularization	0.9856	0.9702
250HLN+L2 regularization	0.9465	0.9453

The number of epochs is set to 10 and the learning rate to 0.01 for all configurations.

The table shows that the regularization term has a negative impact on the training accuracy when compared to the MLPs with no regularization term. However, the testing error is now approximately the same as the training error, meaning that the model extrapolates well to unseen data. This is opposed to the case with no regularization, where the testing accuracy is lower than the training accuracy.

[0]: